End-to-end network programmability

From core switches to end hosts

Gregorio Procissi   gregorio.procissi@unipi.it
Giuseppe Lettieri   giuseppe.lettieri@unipi.it

Agenda

Network programmability all over the network

- Part I (Gregorio Procissi)
  - From original SDN to programmable network data-plane
  - In-network computing: programmable switches and the P4 language
  - P4 in practice: running a programmable software switch in an emulated network environment

- Part II (Giuseppe Lettieri)
  - End-host computing
    - in-kernel networking with extended Berkeley Packet Filter (eBPF)
  - eBPF in practice
    - programming simple applications on a Linux machine

Part I

Gregorio Procissi

In-network data plane programmability

Data Plane/Control Plane...

- Data Plane (Forwarding plane)
  - Processing of data packets, in particular forwarding

- Control plane
  - Intelligence of the network, defines how to handle packets
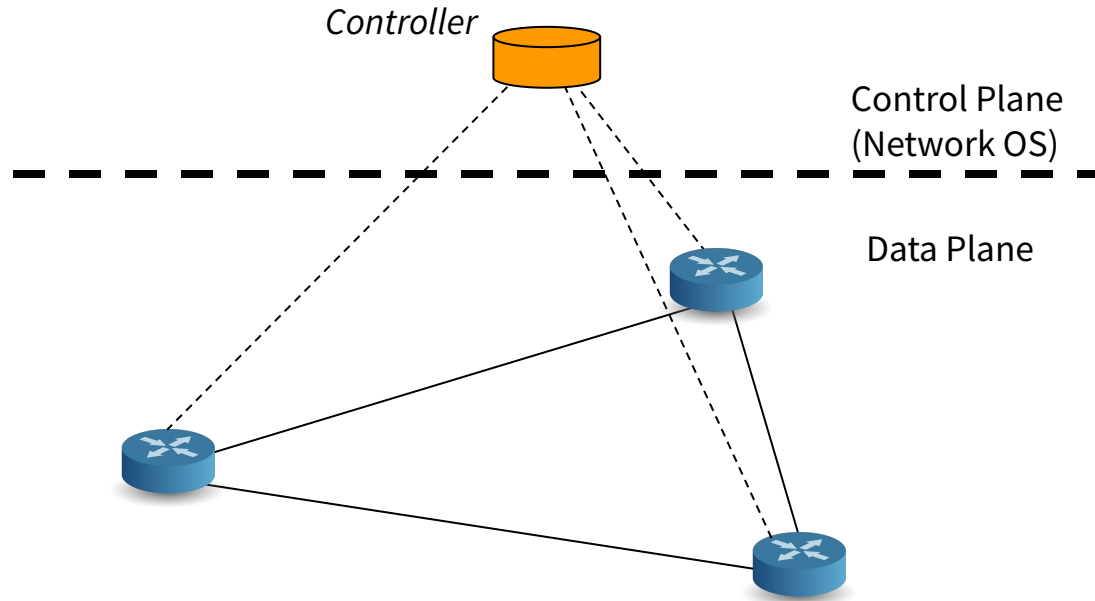 - example: *routing*

... In traditional networks

- Data and control planes vertically integrated
  - in every network node
  - control plane distributed across nodes (routers)
- *Destination based forwarding*
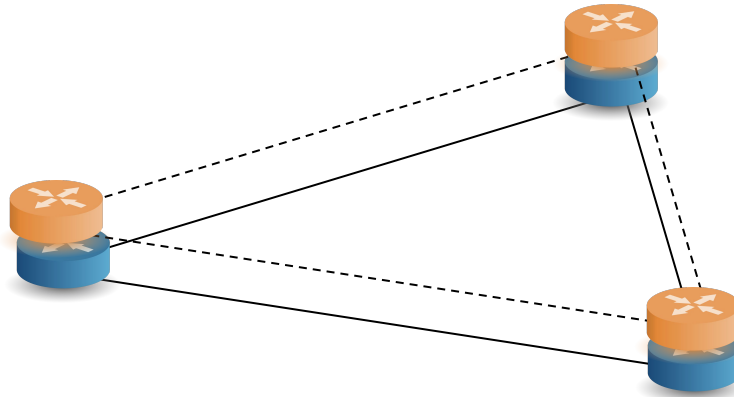
... and in Software Define Networks

- Data plane and control plane are logically separated

Controller

Control Plane
(Network OS)

Data Plane

Software Defined Networking (SDN)

- Flow based forwarding
  - **flow**: set of header fields values used as **matching rules**
  - **actions**: to be applied to packets belonging to the same flow

- Extended concept of forwarding node
  - a switch may be **instructed** to be a router, firewall, load-balancer, and so on

- At this stage network applications run on top of the SDN controller
  - just as processes run on top of operating systems
  - programmability is quite classic
  - does it **scale** at high traffic rates...?

Openflow

- The first successful incarnation of the SDN paradigm, from Stanford University
- from OF paper (2008): *"amenable to high performance and low-cost implementation; capable of supporting a broad range of research; and consistent with vendors' need for closed platforms"*
- Essentially, OpenFlow is an **API** to a switch flow table
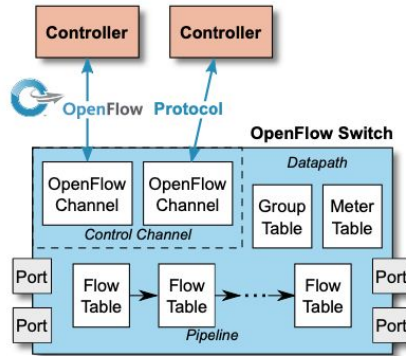
Match-Action Tables

- Match-Action Table
  - Match:
    - binary match (0,1)
    - ternary match (0,1,*)
  - Action:
    - drop, forward, modify, go to another table

Switch Port	Mac src	Mac dst	Eth type	IP src	IP dst	TCP s_port	TCP d_port	Action
*	*	68:5b:35:83:c8:56	*	*	*	*	*	Forward to Port P2

OpenFlow tables (examples)

L2 switch

Switch Port	Mac src	Mac dst	Eth type	IP src	IP dst	TCP s_port	TCP d_port	Action
*	*	68:5b:35:83:c8:56	*	*	*	*	*	Forward to Port P2

L3 forwarding

Switch Port	Mac src	Mac dst	Eth type	IP src	IP dst	TCP s_port	TCP d_port	Action
*	*	*	0x0800	*	131.114.52.*	*	*	Forward to Port P3

L4 Firewall

Switch Port	Mac src	Mac dst	Eth type	IP src	IP dst	TCP s_port	TCP d_port	Action
*	*	*	0x0800	131.114.52.*	*	*	6969	Drop

OpenFlow Switching platforms

- Commodity CPU: software switch (OVS, OFSoftSwitch, etc…)

- FPGA, NetFPGA

- Network Processor (NPU)

- Switching chips
  - ~ 100x faster than CPUs
  - ~ 10x faster than NPU

Is it really data plane programming?

- To add *expressiveness* the protocol has become more and more **complex**
  - 12 fields match on OF 1.0 ---> 41 in OF 1.4
- **Fixed** type tables
- **No custom actions** upon matching
  - only a set of "pre-cooked" actions
  - add brand new headers
  - …
- No room for defining custom **state variables** (across different packets)

```
enum oxm_ofb_match_fields {
  OFPXMT_OFB_IN_PORT,
  OFPXMT_OFB_IN_PHY_PORT,
  OFPXMT_OFB_METADATA,
  OFPXMT_OFB_ETH_DST ,
  OFPXMT_OFB_ETH_SRC ,
  OFPXMT_OFB_ETH_TYPE ,
  OFPXMT_OFB_VLAN_VID ,
  OFPXMT_OFB_VLAN_PCP ,
  OFPXMT_OFB_IP_DSCP ,
  OFPXMT_OFB_IP_ECN ,
  OFPXMT_OFB_IP_PROTO ,
  OFPXMT_OFB_IPV4_SRC ,
  OFPXMT_OFB_IPV4_DST ,
  OFPXMT_OFB_TCP_SRC ,
  OFPXMT_OFB_TCP_DST ,
  OFPXMT_OFB_UDP_SRC ,
  OFPXMT_OFB_UDP_DST ,
  OFPXMT_OFB_SCTP_SRC ,
  OFPXMT_OFB_SCTP_DST ,
  OFPXMT_OFB_ICMPV4_TYPE ,
  OFPXMT_OFB_ICMPV4_CODE ,
  OFPXMT_OFB_ARP_OP ,
  OFPXMT_OFB_ARP_SPA ,
  OFPXMT_OFB_ARP_TPA ,
  OFPXMT_OFB_ARP_SHA ,
  OFPXMT_OFB_ARP_THA ,
  OFPXMT_OFB_IPV6_SRC ,
  OFPXMT_OFB_IPV6_DST ,
  OFPXMT_OFB_IPV6_FLABEL ,
  OFPXMT_OFB_ICMPV6_TYPE ,
  OFPXMT_OFB_ICMPV6_CODE ,
  OFPXMT_OFB_IPV6_ND_TARGET ,
  OFPXMT_OFB_IPV6_ND_SLL ,
  OFPXMT_OFB_IPV6_ND_TLL ,
  OFPXMT_OFB_MPLS_LABEL ,
  OFPXMT_OFB_MPLS_TC ,
  OFPXMT_OFP_MPLS_BOS ,
  OFPXMT_OFB_PBB_ISID ,
  OFPXMT_OFB_TUNNEL_ID ,
  OFPXMT_OFB_IPV6_EXTHDR ,
  OFPXMT_OFB_PBB_UCA
};
```

```
enum ofp_action_type {
  OFPAT_OUTPUT,
  OFPAT_COPY_TTL_OUT,
  OFPAT_COPY_TTL_IN,
  OFPAT_SET_MPLS_TTL,
  OFPAT_DEC_MPLS_TTL,
  OFPAT_PUSH_VLAN,
  OFPAT_POP_VLAN,
  OFPAT_PUSH_MPLS,
  OFPAT_POP_MPLS,
  OFPAT_SET_QUEUE,
  OFPAT_GROUP,
  OFPAT_SET_NW_TTL,
  OFPAT_DEC_NW_TTL,
  OFPAT_SET_FIELD,
  OFPAT_PUSH_PBB,
  OFPAT_POP_PBB,
  OFPAT_EXPERIMENTER
};
```

Deeper network programmability

- Why not defining flexible parsers to match arbitrary header fields?
- Why not defining custom tables for matching rules?
- Why not implementing custom actions to be applied upon matching rules?
- Why not supporting custom stateful operations?
- The answer… again came from Stanford University… and is… **P4** **Programming Protocol-independent Packet Processors + PISA**

Reconfigurable Match Tables (RMTs)

- **Reconfigurable Match Table** (RMT)
  - Set of pipelined stages each with a match table of **arbitrary depth** and **width**
    - ex: IP match: 32-bit depth, 256k width
    - ex: Ethernet: 48-bit addresses width, 64k width
- RMT introduces reconfigurability:
  - field definitions *can be altered* and *new fields* added
  - the number, topology, widths, and depths of match tables *can be specified*, subject only to an overall resource limit on the number of matched bits
  - *new actions may be defined*, such as writing new header fields
  - arbitrarily modified packets can be placed in specified queue(s), for output at any subset of ports, with a queuing discipline specified for each queue
- Configuration managed by an SDN controller

Mac src	Mac dst	Eth type	Action
*	68:5b:35:83:c8:56	*	Goto Table 2
...	...	...	...
*	Aa:bb:00:cc:00:aa	*	Goto Table 2

Table 1

IP src	IP dst	Action
*	*	Forward to Port P2
*	131.114.52.*	Forward to Port P3
131.114.52.*	*	Drop

Table 2

PISA logical architecture

Protocol Independent Switch Architecture (PISA)

- **Programmable** *parser*
    - yields a **packet header vector**
        - header fields + metadata
- **Programmable** *match-action pipeline*
    - Packet header vectors flows through a sequence of logical match stages that run in series or in parallel
- **Programmable** *deparser*
    - Recompose the packet be serializing headers in the desired order
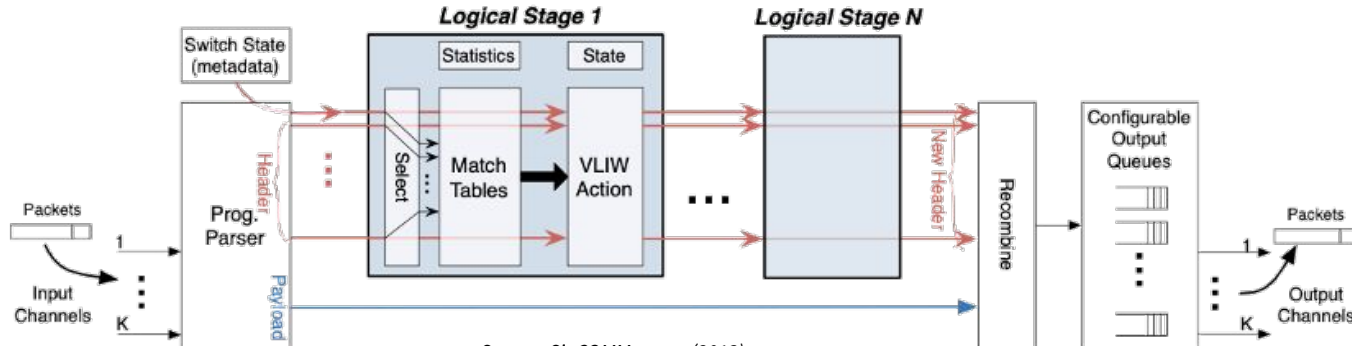
Source: SigCOMM paper (2013)

PISA: Protocol-Independent Switch Architecture

Programmer declares the headers that should be recognized and their order in the packet

Programmer defines the tables and the exact processing algorithm

Programmer declares how the output packet will look on the wire

Programmable Match-Action Pipeline

Programmable Parser

Programmable Deparser
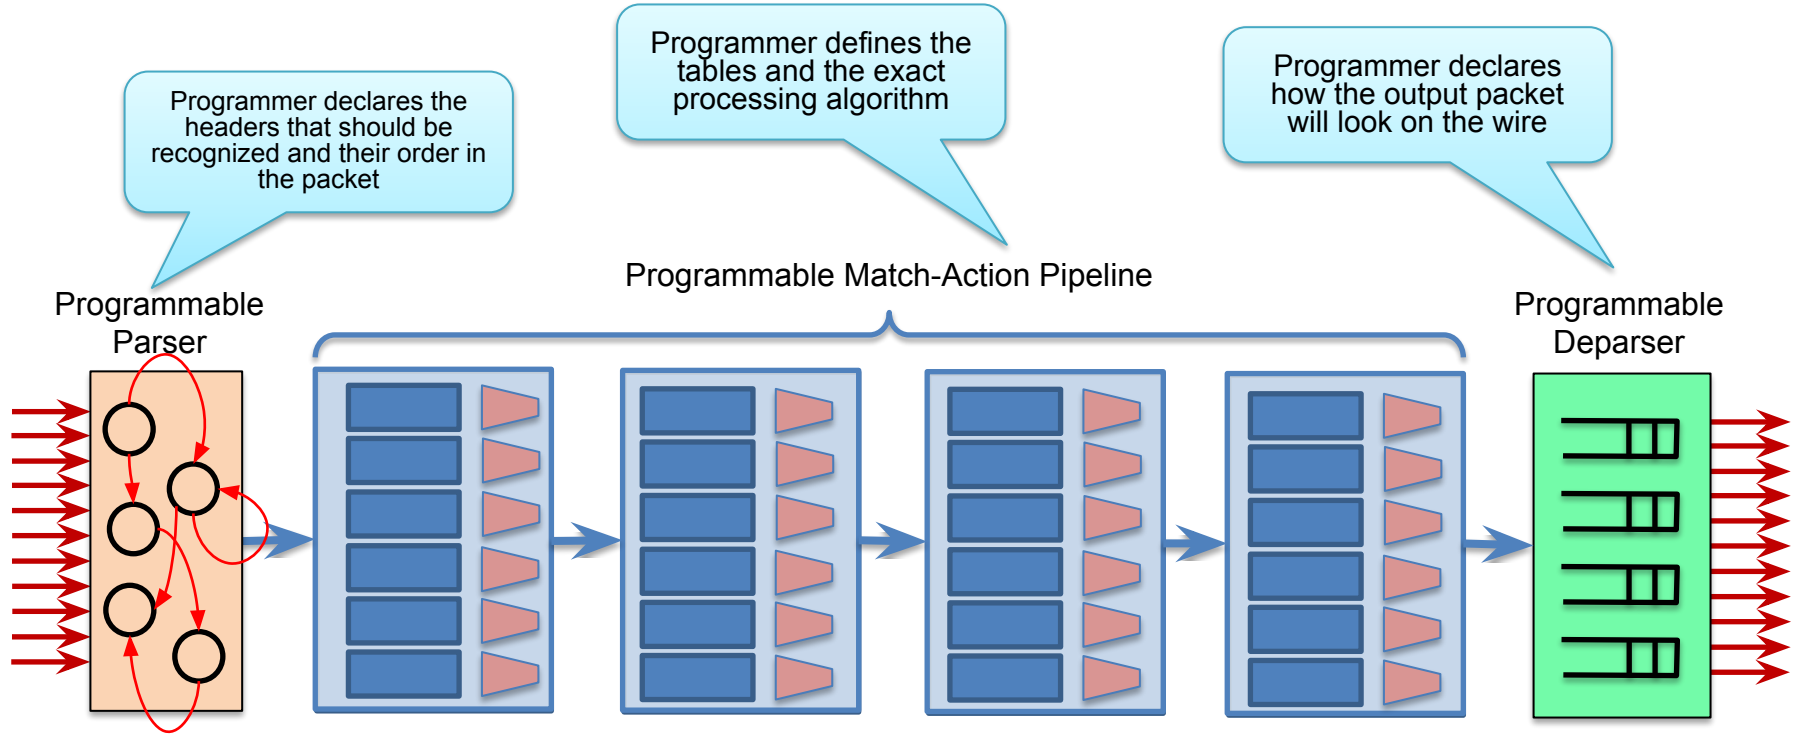
Source: p4.org

Packet Processing in PISA

Source: p4.org
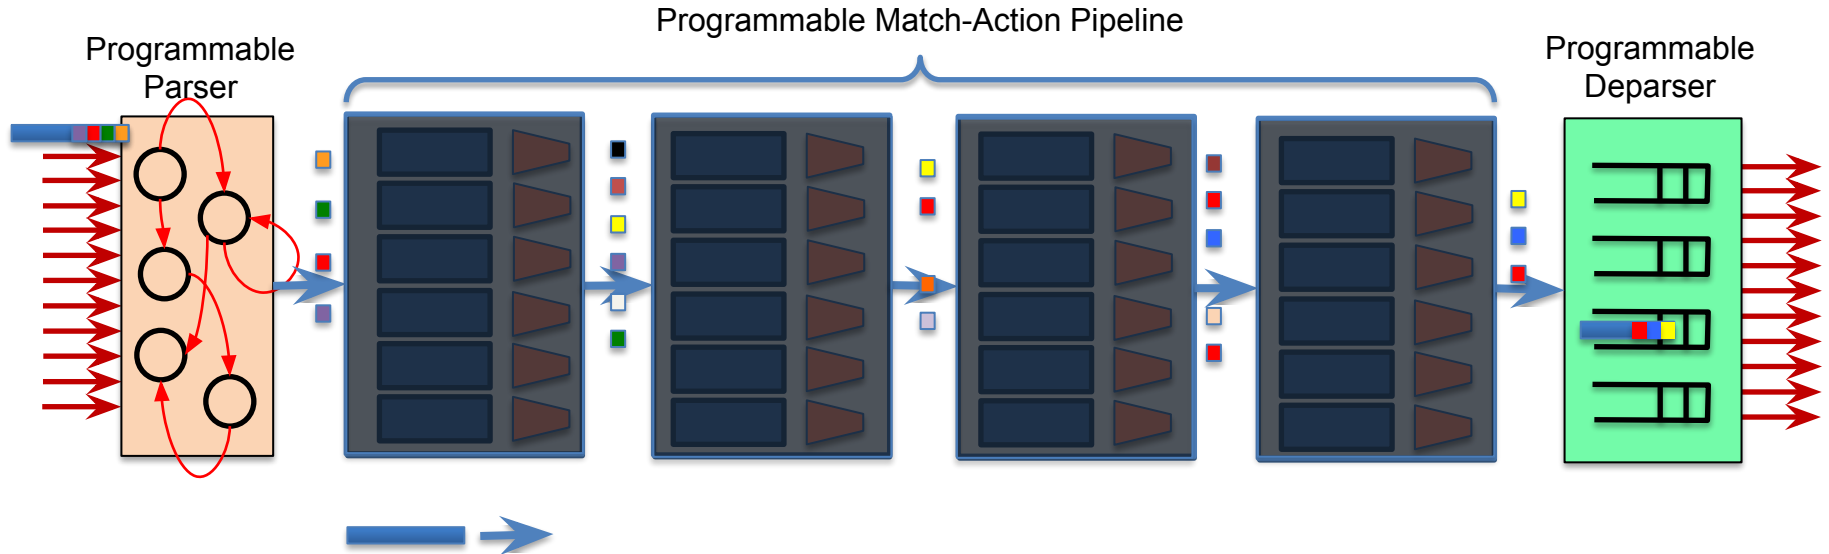
P4 architectures

The architecture is what the programmer sees, a logical view of the processing hide the underlying hw details to the programmer

my_program.p4
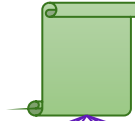
Written against a specific architecture
Defines the processing of each block

Packet Queuing, Replication & Scheduling
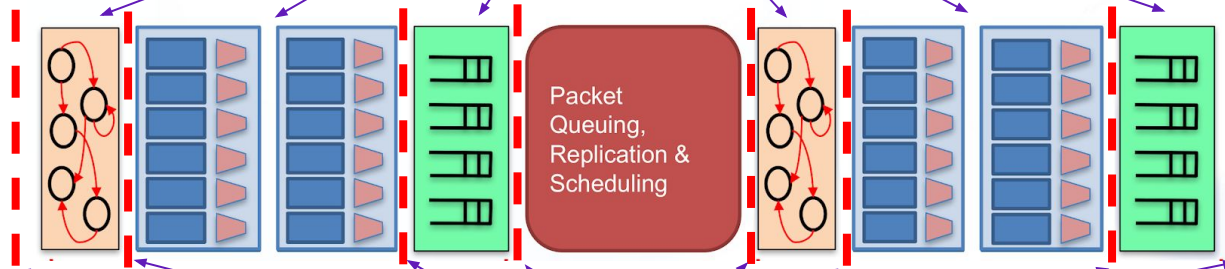
architecture.p4
Provided by switch vendor
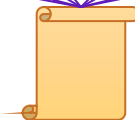
Defines which blocks are available, the interfaces of each block, and their capabilities

Source: Open Networking Foundation

20

Architectures and targets
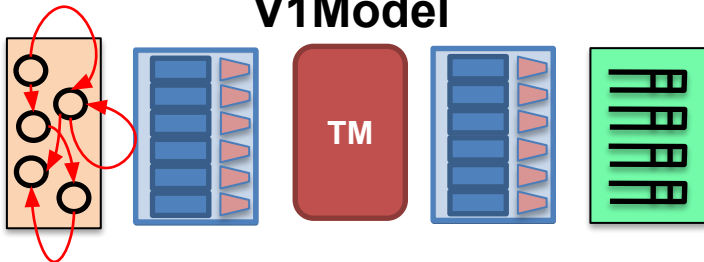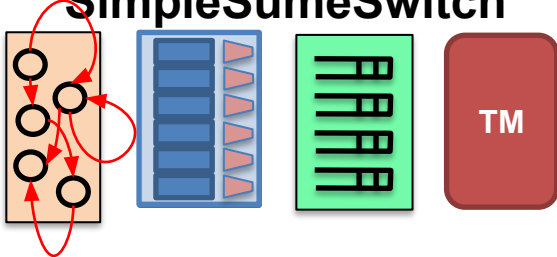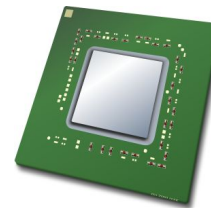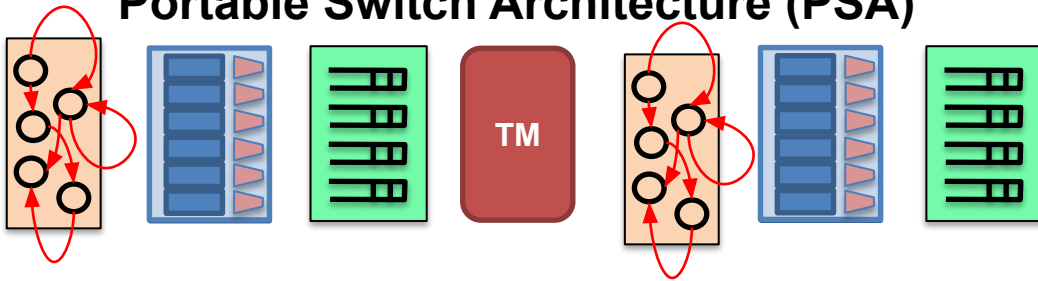
Source: p4.org

The v1Model architecture for the BMv2 soft switch

- Implemented in Bmv2's simple_switch target (not the only possible!!)
- Very similar to the original PISA architecture

Parser   **Checksum Verification / Ingress Match-Action**   **Checksum Update / Egress Match-Action**   **Deparser**

Traffic Manager
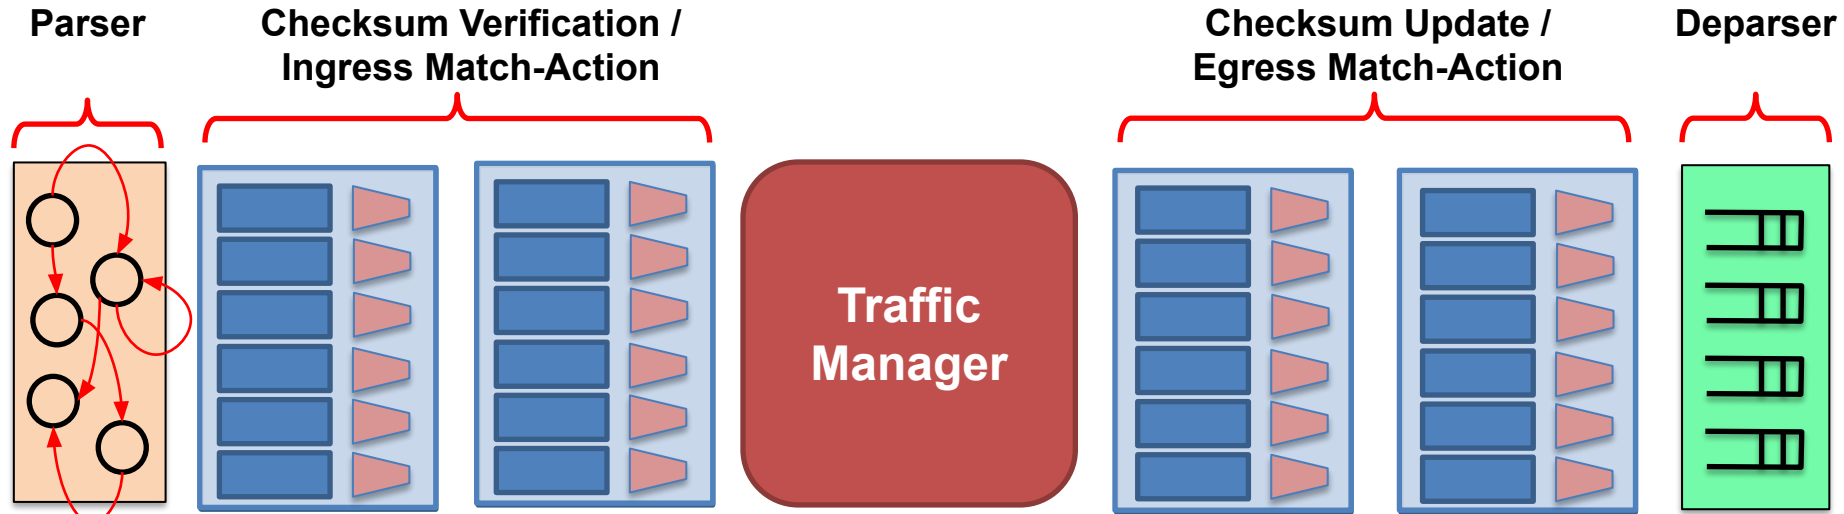
Source: p4.org

P4: programming a target

Vendor supplied
User supplied
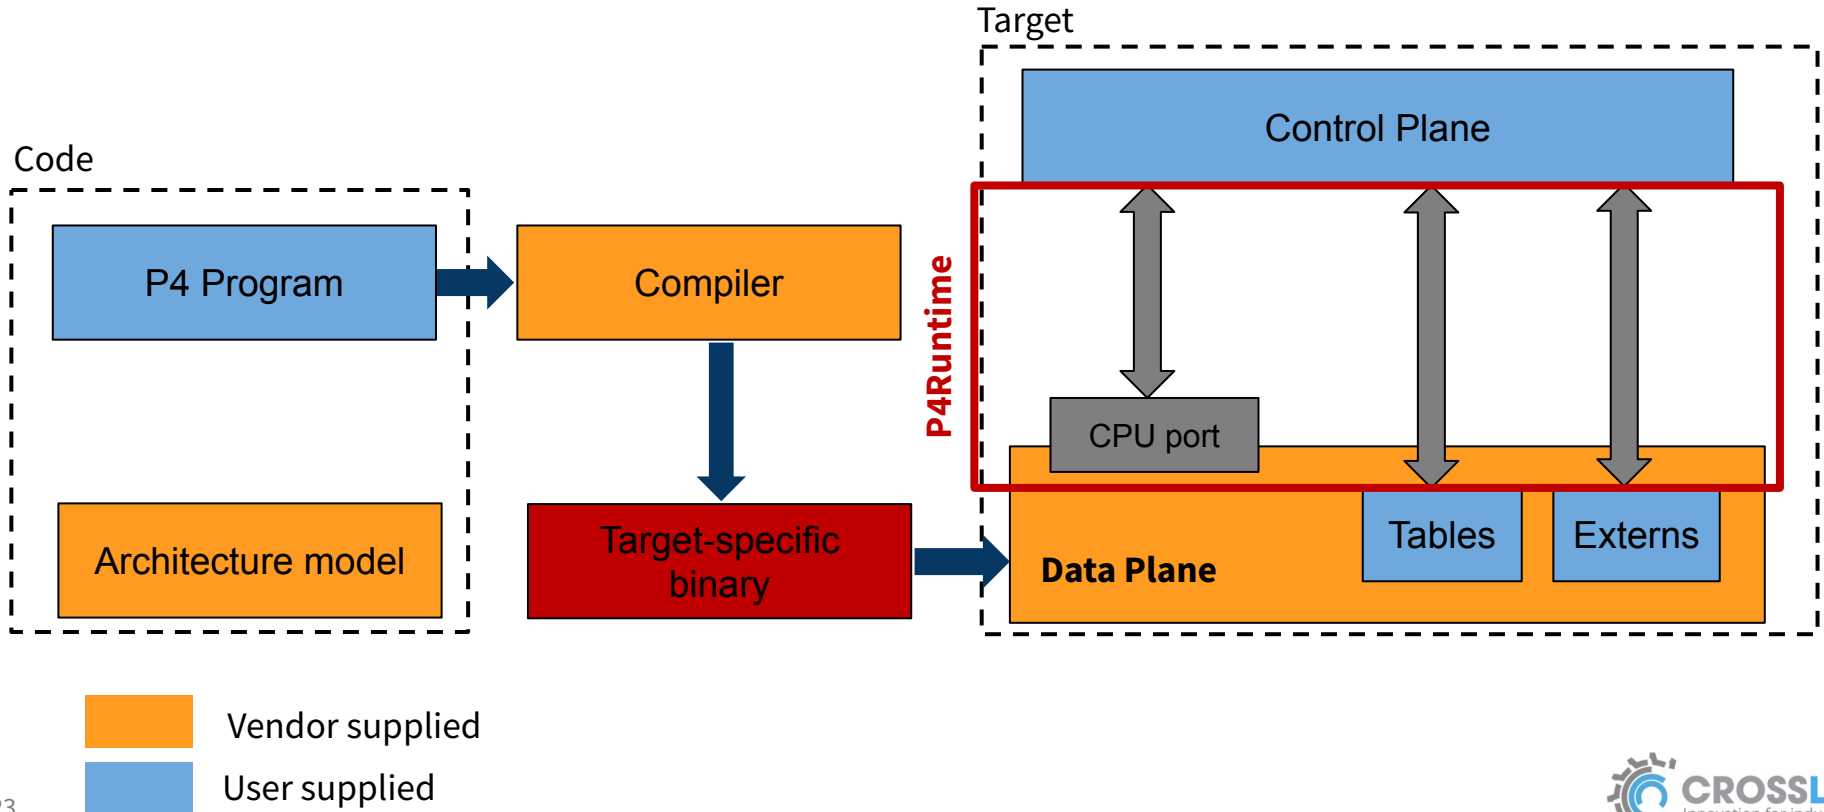
A real running testbed?

- Not very practical…
  - Need a real P4 switch and too many PCs…

Controller

P4 switch
connected to hosts
and controller
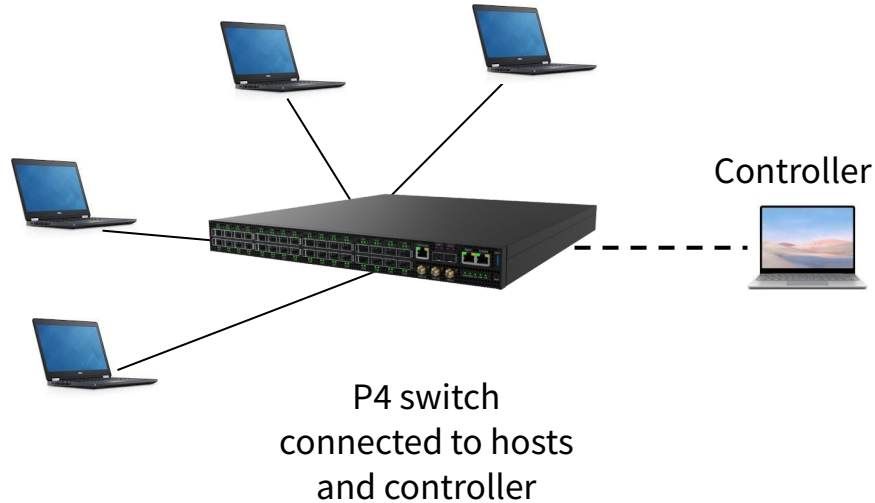
A real running testbed?

- Again, not very practical…
  - Have not many laptops…

Controller

P4Runtime server

Linux server with
programmable
soft switch
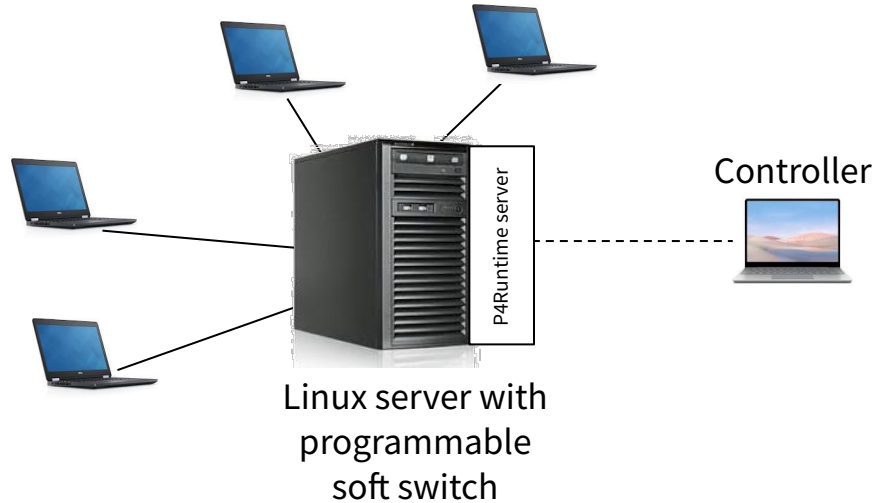
P4 (virtual) running environment
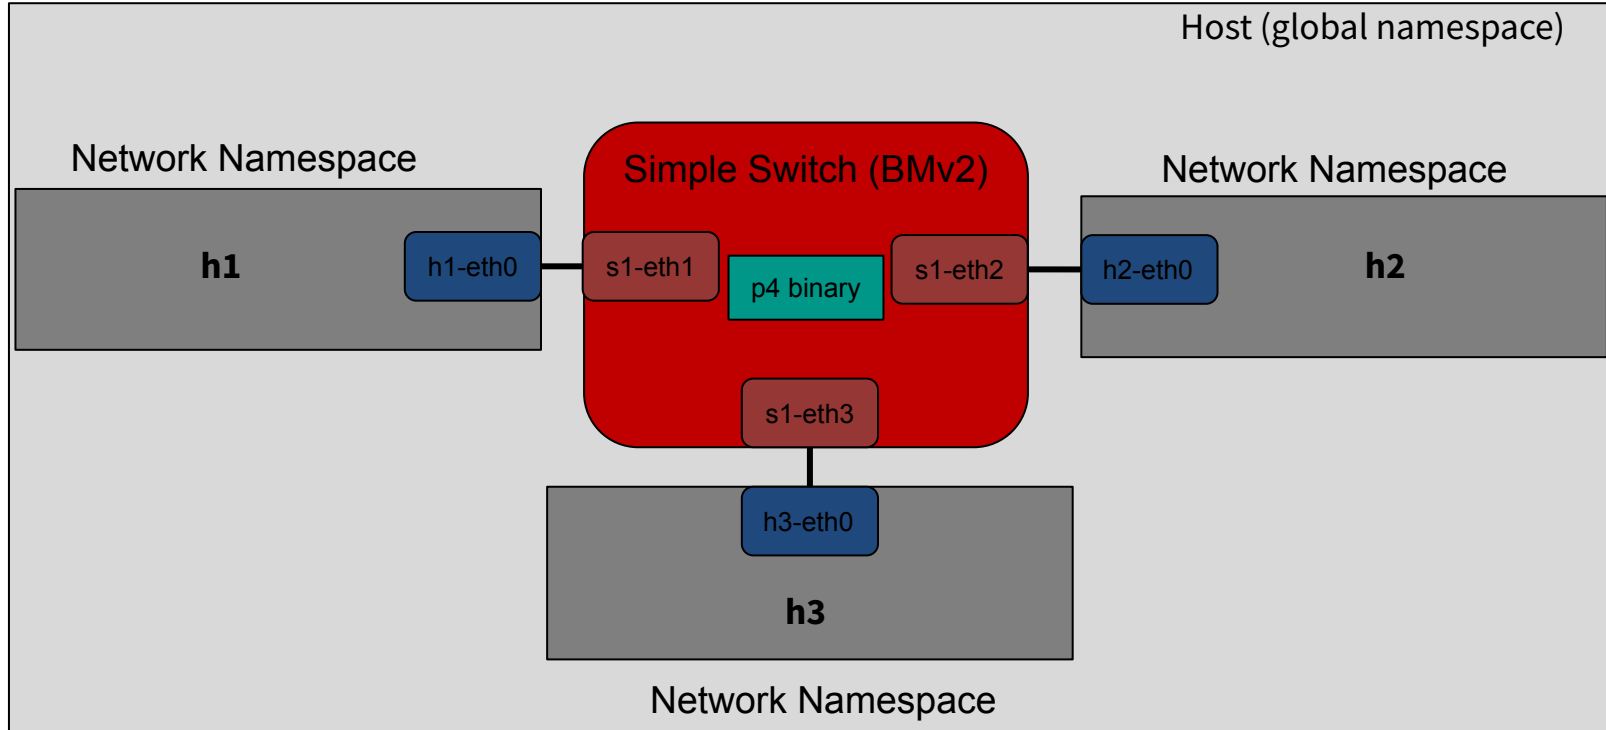
What we need

- A Linux machine with:
  - The **BMv2** software switch
  - **p4c**: the reference P4 compiler
  - **Mininet**: a lightweight network emulation environment
- As lot of dependencies need to be installed, I suggest downloading a VM with everything in it:
  - Tutorial from P4 official page:
    - **https://github.com/p4lang/tutorials**
    - lots of exercises in there
  - github repository of examples presented today
    - **https://github.com/grp-xx/2023-5G-SummerSchool**

p4 program structure

```
#include <core.p4>
#include <v1model.p4>
```
← Libraries

```
struct metadata {
}
struct headers {
}
```
Declarations

```
parser MyParser(packet_in packet,
                out headers hdr,
                inout metadata meta,
                inout standard_metadata_t standard_metadata) {

    state start{
        transition accept;
    }
}
```
← Parser

```
control MyVerifyChecksum(inout headers hdr, inout metadata meta) {
    apply {  }
}
```

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {

    action forward(bit<9> egress_port){
        standard_metadata.egress_spec = egress_port;
    }

    table repeater {...}

    apply {
        repeater.apply();
    }
}
```
← Control for Ingress Pipeline

Egress Pipeline →
```
control MyEgress(inout headers hdr,
                 inout metadata meta,
                 inout standard_metadata_t standard_metadata) {
    apply {  }
}
```

```
control MyComputeChecksum(inout headers  hdr, inout metadata meta) {
    apply { }
}
```

Packet re-assembly →
```
control MyDeparser(packet_out packet, in headers hdr) {
    apply {

    /* Deparser not needed */

    }
}
```

```
V1Switch(
    MyParser(),
    MyVerifyChecksum(),
    MyIngress(),
    MyEgress(),
    MyComputeChecksum(),
    MyDeparser()
) main;
```
← main

29

v1Model standard_metadata

```
struct standard_metadata_t {
    bit<9>  ingress_port;
    bit<9>  egress_spec;
    bit<9>  egress_port;
    bit<32> clone_spec;
    bit<32> instance_type;
    bit<1>  drop;
    bit<16> recirculate_port;
    bit<32> packet_length;
    bit<32> enq_timestamp;
    bit<19> enq_qdepth;
    bit<32> deq_timedelta;
    bit<19> deq_qdepth;
    bit<48> ingress_global_timestamp;
    bit<32> lf_field_list;
    bit<16> mcast_grp;
    bit<1>  resubmit_flag;
    bit<16> egress_rid;
    bit<1>  checksum_error;
}
```

The arrival port of the packet

The port the packet should be sent to

The departure port (can be read in the egress pipeline only)

P4 Types

- Statically typed language with
  - base types
  - composed types

```
bool
bit<W>
int<W>
varbit<W>
enum
typedef bit<48> macAddr_t;
...
no float, no strings!
```

```
struct standard_metadata_t {
  bit<9> ingress_port;
  bit<9> egress_spec;
  bit<9> egress_port;
... }
```

```
header Ethernet_h {
  bit<48> dstAddr;
  bit<48> srcAddr;
  bit<16> etherType;
}
Ethernet_h ethernet;
```

```
tuple<bit<32>, bool> x;
x = { 10, false };
```

```
struct headers{
  Ethernet_h ethernet;
  IPv4_h     ipv4;
  ... }
headers hdr;
```

```
header Mpls_h {
  bit<20> label;
  bit<3>  tc;
  bit     bos;
  bit<8>  ttl;
}
```

```
Mpls_h[10] mpls; //stack
```

```
header_union IP_h {
  IPv4_h v4;
  IPv6_h v6;
}
IP_h ipHeader;
```

P4 operations

- Arithmetic
  - +, -, *   / * no division!!! */
- Logical
  - ~, &, |, ^, <<, >>
- Bit-slicing [m:l]
- Bit concatenation ++
- **No modulo operation!!!**

P4 Variables and constants

- Variables
  - Have local scope
  - not maintained upon the next invocation
  - **cannot be used to save states**!
    - use tables and extern for that

```
bit<8> x = 10;

typedef bit<16> TcpPort;
TcpPort s_port;
s_port = 10000;
```

- Constants

```
const bit<8> x = 10;

typedef bit<16> TcpPort;
const TcpPort s_port = 10000;
```

P4 main statements

- **return**
  - terminate action or control in which it is contained
- **exit**
  - terminate of all blocks in execution
- **conditions**
  - cannot be used in parsers

  ```
  if (x==100) {...} else {...}
  ```

- **switch**
  - can be used in control blocks only

  ```
  switch (hdr.ethernet.etherType) {
      0x86dd: { /* body omitted */ }
      0x0800: { /* body omitted */ }
      0x0802: { /* body omitted */ }
      0xcafe: { /* body omitted */ }
      default: { /* body omitted */ }
  }
  ```

- **Forget about loops in P4!!!**

P4 parser

- Finite State Machine
  - states
    - predefined **start**, **accept**, **reject**
    - + user defined ones
  - transitions
  - loops are allowed (e.g. tunneling)
  - more advanced methods: `verify,lookahead,...`
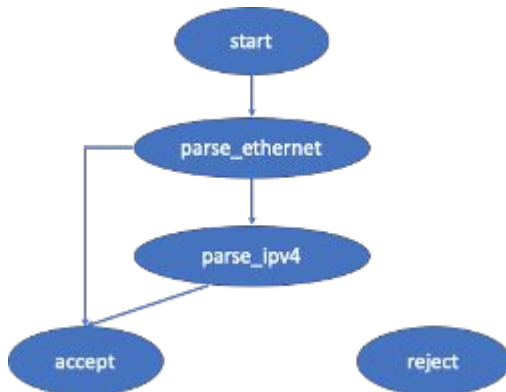
```
parser MyParser(packet_in packet,
          out headers hdr,
          inout metadata meta,
          inout standard_metadata_t std_meta) {

state start {
  transition parse_ethernet;
}

state parse_ethernet {
  packet.extract(hdr.ethernet);
  transition select(hdr.ethernet.etherType)
{
    0x800: parse_ipv4;
    default: accept;
  }
}

state parse_ipv4 {
...
}
}
```

P4 control

- Control blocks implements:
    - match actions pipelines
    - deparser
    - additional packet processing (e.g. checksum update)
- Control blocks contain:
    - Tables
        - match a key and return an action (and data)
    - Actions
        - pretty much like C functions, for software re-use
        - no return value, though
        - parameters have **directions**
            - `in`   ReadOnly in the action
            - `out`   to be written in the action
            - `inout` ReadWrite
    - Control flow
        - blocks of imperative code (without loops!)
        - may contain advanced concepts such as cloning packets, sending packets to control plane, etc.

P4 action and control flow

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t std_meta) {
  /* Declarations region */
  bit<48> tmp;

  apply {
    /* Control Flow */
    tmp = hdr.ethernet.dstAddr;
    hdr.ethernet.dstAddr = hdr.ethernet.srcAddr;
    hdr.ethernet.srcAddr = tmp;
    std_meta.egress_spec = std_meta.ingress_port;
  }
}
```

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t std_meta) {

  action swap_mac(inout bit<48> src,
                  inout bit<48> dst) {
    bit<48> tmp = src;
    src = dst;
    dst = tmp;
  }

  apply {
    swap_mac(hdr.ethernet.srcAddr,
             hdr.ethernet.dstAddr);
    std_meta.egress_spec = std_meta.ingress_port;
  }
}
```

P4 Tables

Control plane

key

Match — key

Action — ID — Data

Hit

Default

Headers and metadata

Headers and metadata

ID

Action Code

Data
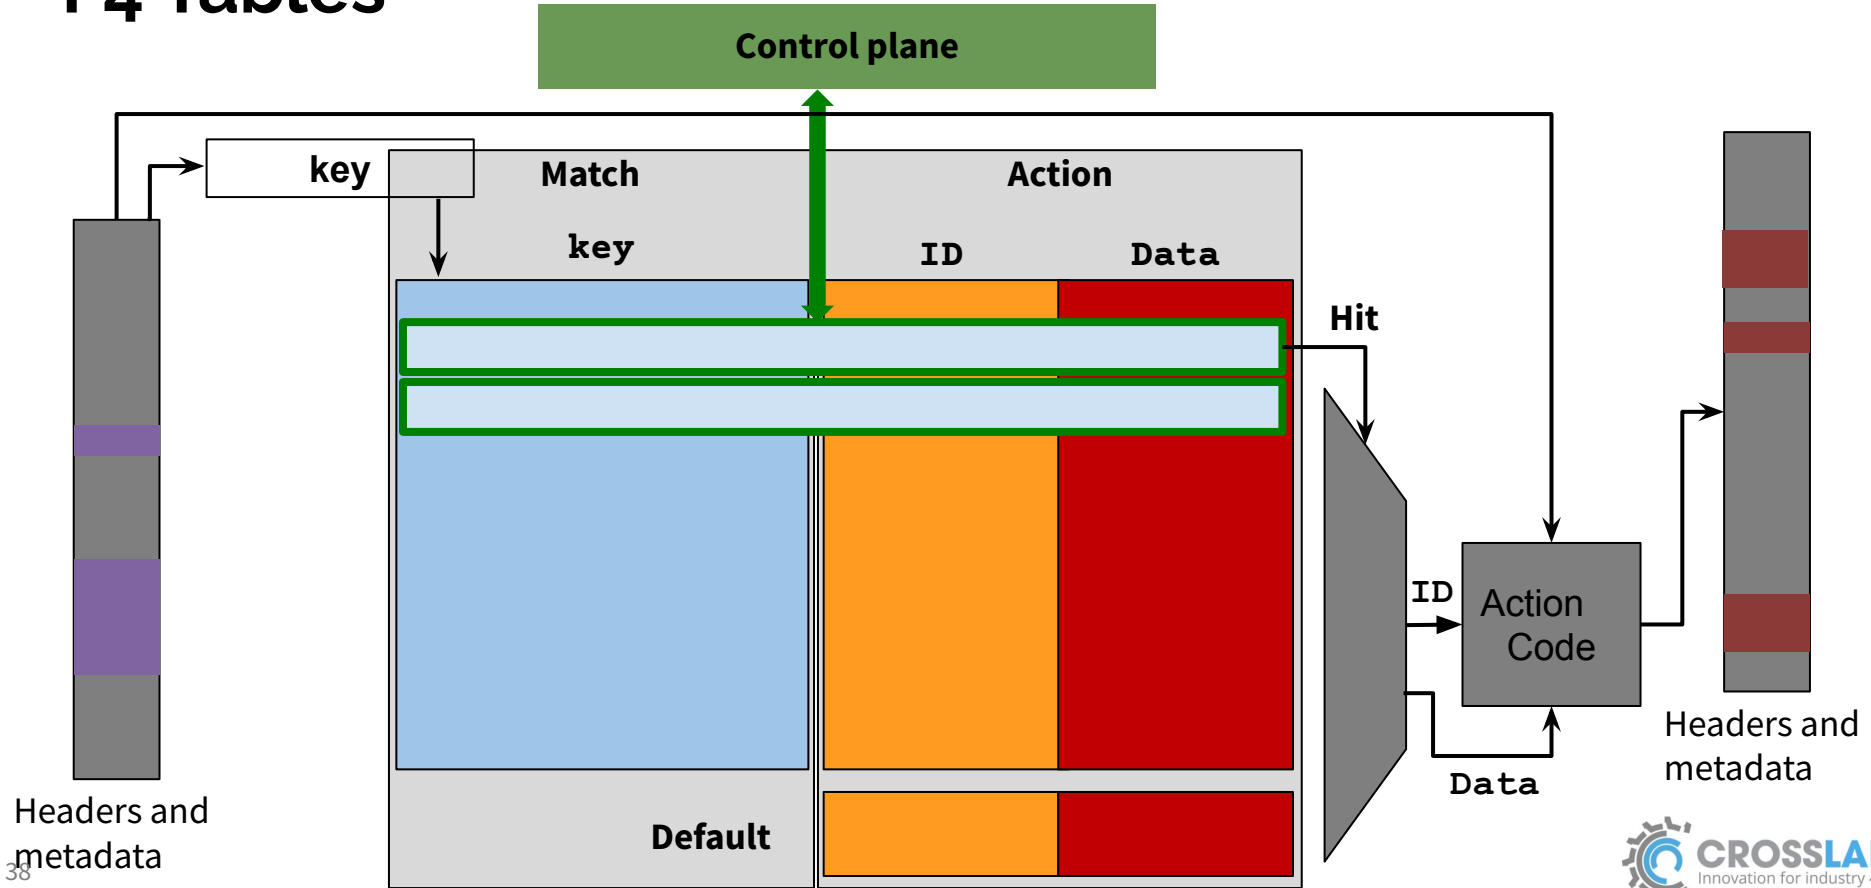
P4 tables in the code...

Table Name

```
table ipv4_lpm {
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    actions = {
        ipv4_forward;
        drop;
        NoAction;
    }

    size = 1024;

    default_action = NoAction;
}
```

Match type

Possible Actions

Max # of table entries

Default action

```
/* core.p4 */
match_kind {
    exact,
    ternary,
    lpm
}


/* v1model.p4 */
match_kind {
    range,
    selector
}


/* Some other architecture */
match_kind {
    ...
}
```

Tables populate
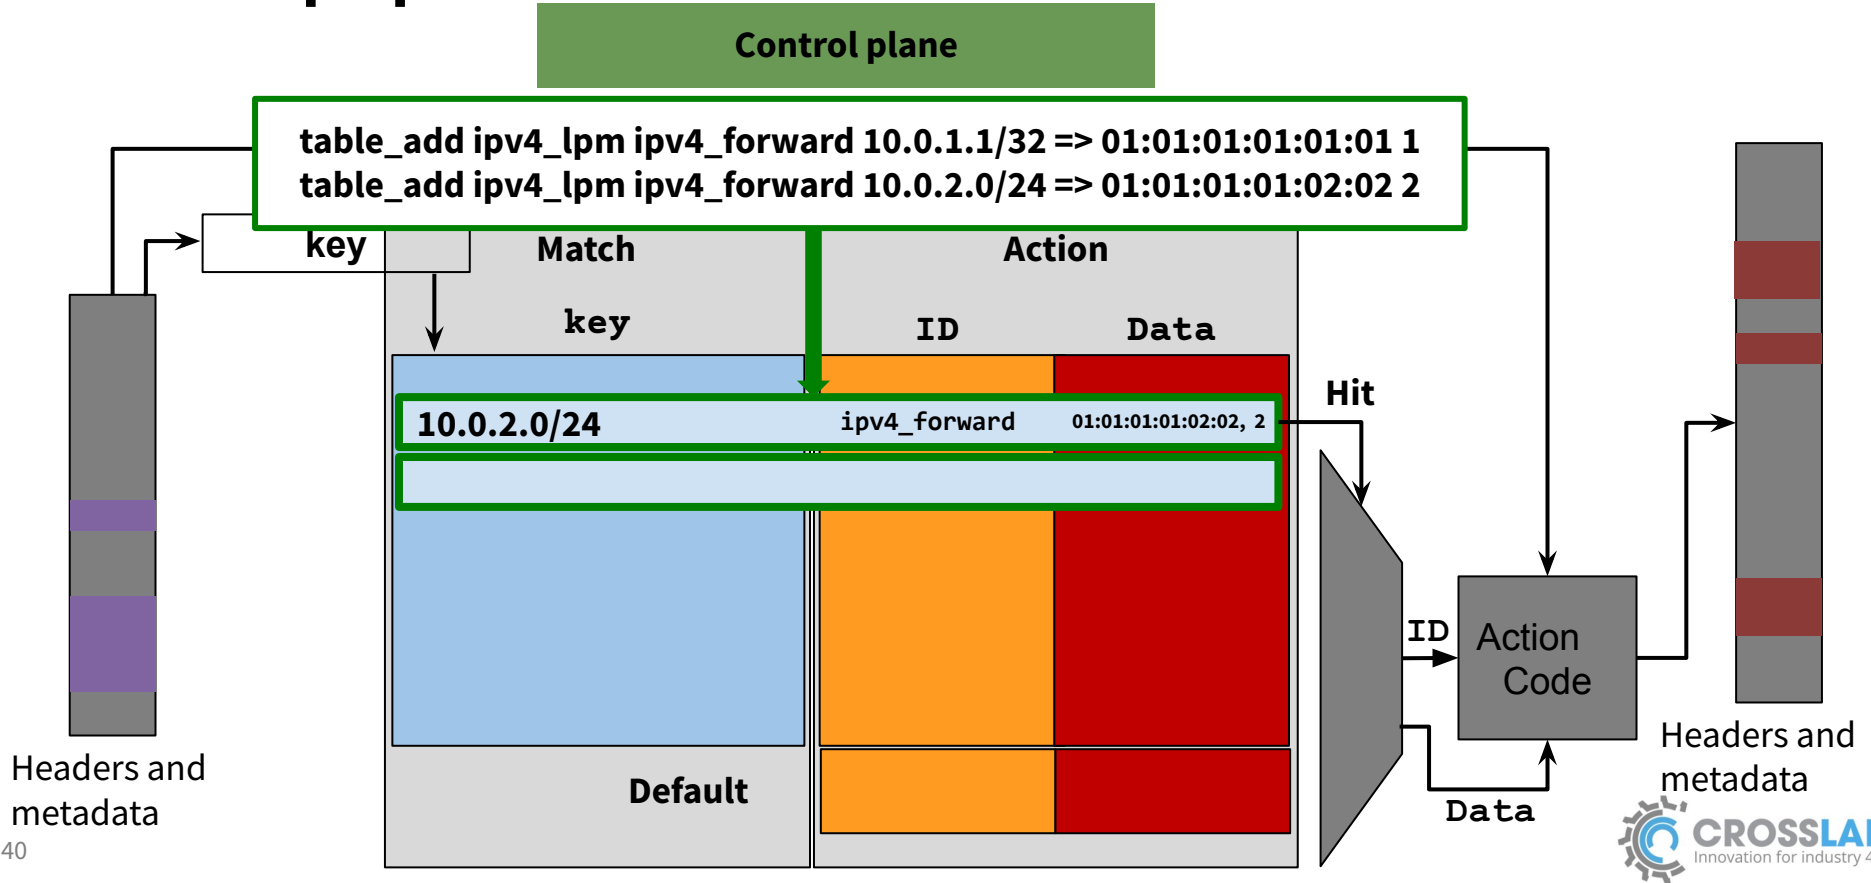
P4 Deparser

- Serializes headers in the desired order and rebuild the packet

```
control DeparserImpl(packet_out packet,
                     in headers hdr) {
  apply {
    ...
    packet.emit(hdr.ethernet);
    packet.emit(hdr.ipv4);
    packet.emit(hdr.tcp);
    packet.emit(hdr.my_own_defined_hdr);

    ...
  }
}
```
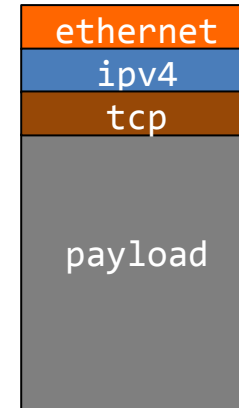
Checksum - Validation and update (1)

- Extern defined in the `v1model`

```
extern void verify_checksum<T, O>( in bool condition,
                                   in T data,
                                   inout O checksum,
                                   HashAlgorithm algo
                                 );
extern void update_checksum<T, O>( in bool condition,
                                   in T data,
                                   inout O checksum,
                                   HashAlgorithm algo
                                 );
```

Checksum - Validation and update (2)

```
control MyComputeChecksum(...) {
  apply {
    update_checksum(
    hdr.ipv4.isValid(),
    {     hdr.ipv4.version,
          hdr.ipv4.ihl,
          hdr.ipv4.diffserv,
          hdr.ipv4.totalLen,
          hdr.ipv4.identification,
          hdr.ipv4.flags,
          hdr.ipv4.fragOffset,
          hdr.ipv4.ttl,
          hdr.ipv4.protocol,
          hdr.ipv4.srcAddr,
          hdr.ipv4.dstAddr },
          hdr.ipv4.hdrChecksum,
    HashAlgorithm.csum16);
  }
}
```

P4 stateful programming

- Stateless objects
  - variables
  - headers
- Stateful objects
  - tables
    - can only be populated by the control plane!
  - Registers ⎫
  - Counters ⎬ *externs*
  - Meters   ⎭
    - used to limit traffic rate

Registers

- Good to store arbitrary data (of type `<Type>`)
- **read** and **write** methods available
- assigned in arrays of length N

`register<Type>(N) reg;`

`<Type>`

```
0
1
2     reg.write(2,value)
3
4
5     reg.read(result,6)
6
7
```

N
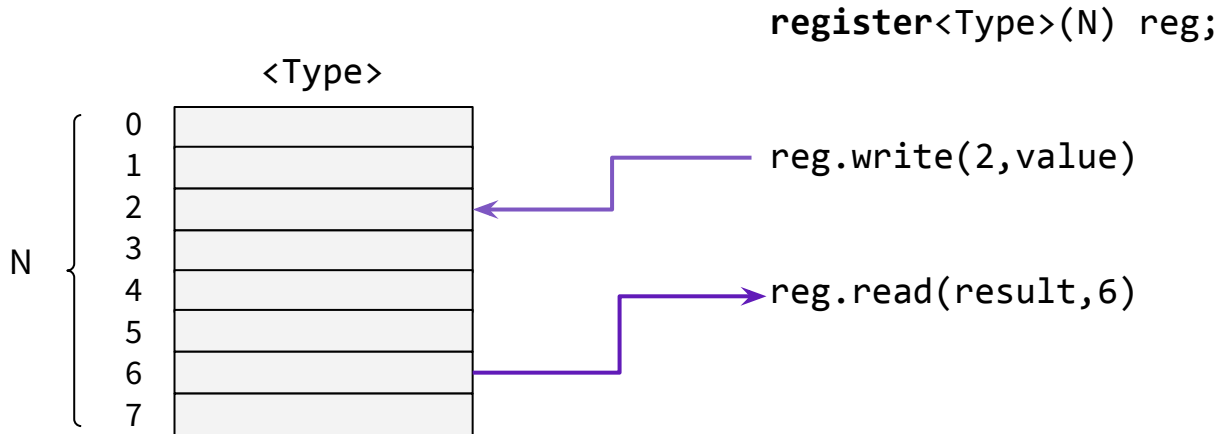
Counters

- Good for... counting ;-)
- Three types only:
  - `packets, bytes, packets_and_bytes`
- Only `count` method available (read available from the control plane)
- Assigned in arrays of length N

`counter(N, <Type>) c;`

`c.count(6)`
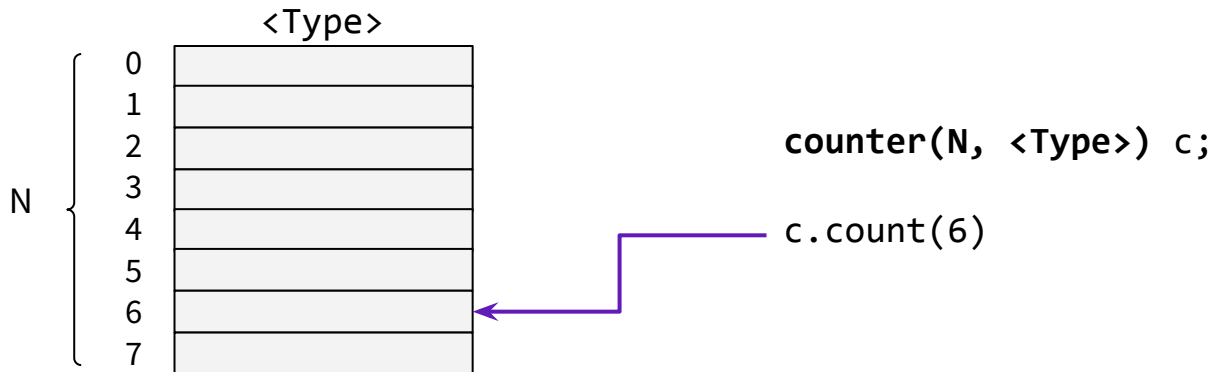
- there are also *direct counters*: special type of counters attached to tables

P4 Hello World: packet reflector

- Not easy to define a hello world for packet processing…
  - can't print a string on video ;-)
- Packet reflector
  - the switch receives the ethernet frame and sends it back to the host (… swaps mac addresses)
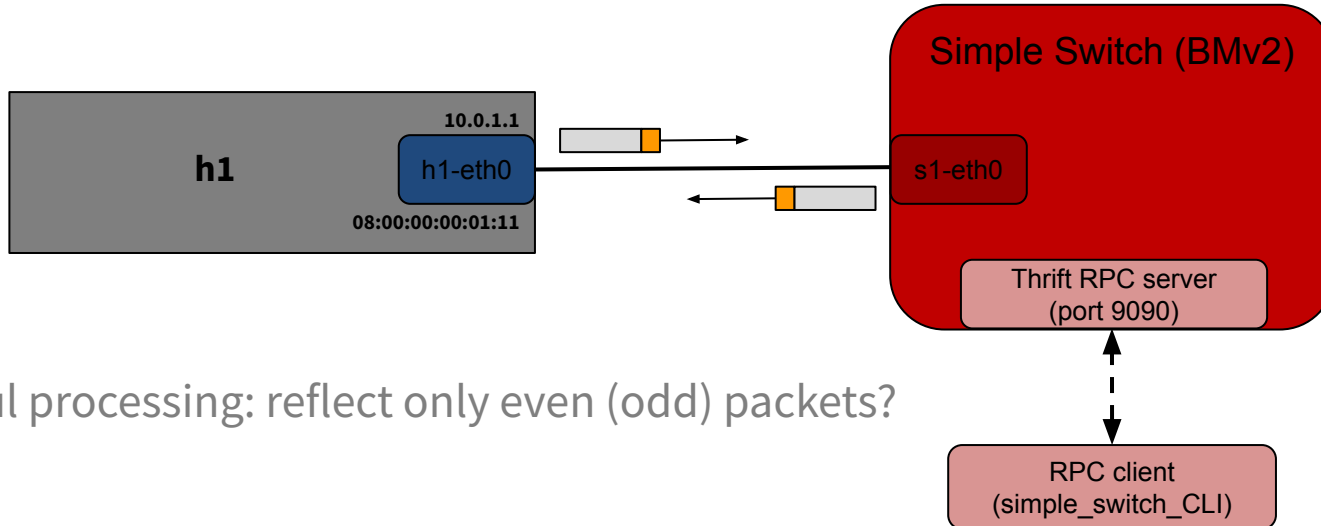
- Stateful processing: reflect only even (odd) packets?

Packet reflector implementation (1)

```
/* -*- P4_16 -*- */
#include <core.p4>
#include <v1model.p4>

/*** H E A D E R S ****/

typedef bit<48> macAddr_t;

header ethernet_t {
    macAddr_t dstAddr;
    macAddr_t srcAddr;
    bit<16> etherType;
}

struct metadata {
    /* empty */
}

struct headers {
        ethernet_t ethernet;
}
```



```
/*** P A R S E R ****/

parser MyParser(packet_in packet,
                out headers hdr,
                inout metadata meta,
                inout standard_metadata_t standard_metadata) {

    state start{
        packet.extract(hdr.ethernet);
        transition accept;
    }
}
```

Packet reflector implementation (2)

```
/*** I N G R E S S   P R O C E S S I N G ****/

control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {


    action swap_macs(inout macAddr_t src, inout macAddr_t dst) {
        macAddr_t tmp = src;
        src = dst;
        dst = tmp;
    }


    apply {
        // Swap MAC addresses.
        swap_macs(hdr.ethernet.srcAddr, hdr.ethernet.dstAddr);

        //Set Output port == Input port
        standard_metadata.egress_spec = standard_metadata.ingress_port;
    }
}
```

```
/*** D E P A R S E R ****/

control MyDeparser(packet_out packet,
                   in headers hdr) {
    apply {
        packet.emit(hdr.ethernet);
    }
}
```

```
/*** S W I T C H  ****/

V1Switch(
    MyParser(),
    MyVerifyChecksum(),
    MyIngress(),
    MyEgress(),
    MyComputeChecksum(),
    MyDeparser()
) main;
```

Reflect one out of two? Count reflected frames?

```
/*** I N G R E S S    P R O C E S S I N G ****/

control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {

    action swap_macs(...){  }

    register<bit<1>>(8) reg;
    apply {
        bit<1> flag;
        bit<32> input_port = (bit<32>) standard_metadata.ingress_port;
        reg.read(flag,input_port);
        reg.write(input_port,flag+1);

        if (flag == 1) {
            mark_to_drop(standard_metadata);
        }
        else {
             swap_macs(hdr.ethernet.srcAddr, hdr.ethernet.dstAddr);
             standard_metadata.egress_spec = standard_metadata.ingress_port;
        }
    }

}
```

```
/*** E G R E S S ****/

control MyEgress(inout headers hdr,
                 inout metadata meta,
                 inout standard_metadata_t standard_metadata)
{

    // Counting the number of reflected frames
    counter(8,CounterType.packets) egress_port_counter;

    apply {
        egress_port_counter.count(
(bit<32>)standard_metadata.egress_port);
    }

}
```

- How do I access **registers** and **counters** from the **control plane**?

```
simple_switch_CLI --thrift-port 9090
> register_read...
> counter_read...
```

And finally using tables...

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t
                                  standard_metadata) {

    register<bit<1>>(8) reg;
    bit<1> flag;

    action swap_macs(inout macAddr_t src, inout macAddr_t dst)
    {
        macAddr_t tmp = src;
        src = dst;
        dst = tmp;
    }

    action frame_reflect() {
        swap_macs(hdr.ethernet.srcAddr, hdr.ethernet.dstAddr);
        standard_metadata.egress_spec =
                          standard_metadata.ingress_port;
    }

    action drop() {
        mark_to_drop(standard_metadata);
    }
```

```
    table odd_even {
    key = {
        flag: exact;
    }

    actions = {
        frame_reflect;
        drop;
    }

    size = 8;
    default_action = drop;
  }

  apply {
    bit<32> input_port = (bit<32>)
                  standard_metadata.ingress_port;
    reg.read(flag,input_port);
    reg.write(input_port,flag+1);
    odd_even.apply();
  }
}
```

- To populate the table...
```
simple_switch_CLI --thrift-port 9090
> table_add odd_even frame_reflect 0 =>
```

Add your own header...

A simple (and useless) example: add the input port number after the ethernet header

- In the declaration section

```
header extra_t {
    bit<16> in_port;
}

struct headers {
    ethernet_t ethernet;
    extra_t extra;
}
```

- In the MyIngress control block

```
  hdr.extra.setValid();
  hdr.extra.in_port = (bit<16>) std_meta.ingress_port;
```

- In the deparser

```
control MyDeparser(packet_out packet, in headers hdr) {

 apply {
        packet.emit(hdr.ethernet);
        packet.emit(hdr.extra);
        }
}
```

L2 Switching

- Switching table provided through the control plane
- Table populated through the simple_switch_CLI (gRPC client)
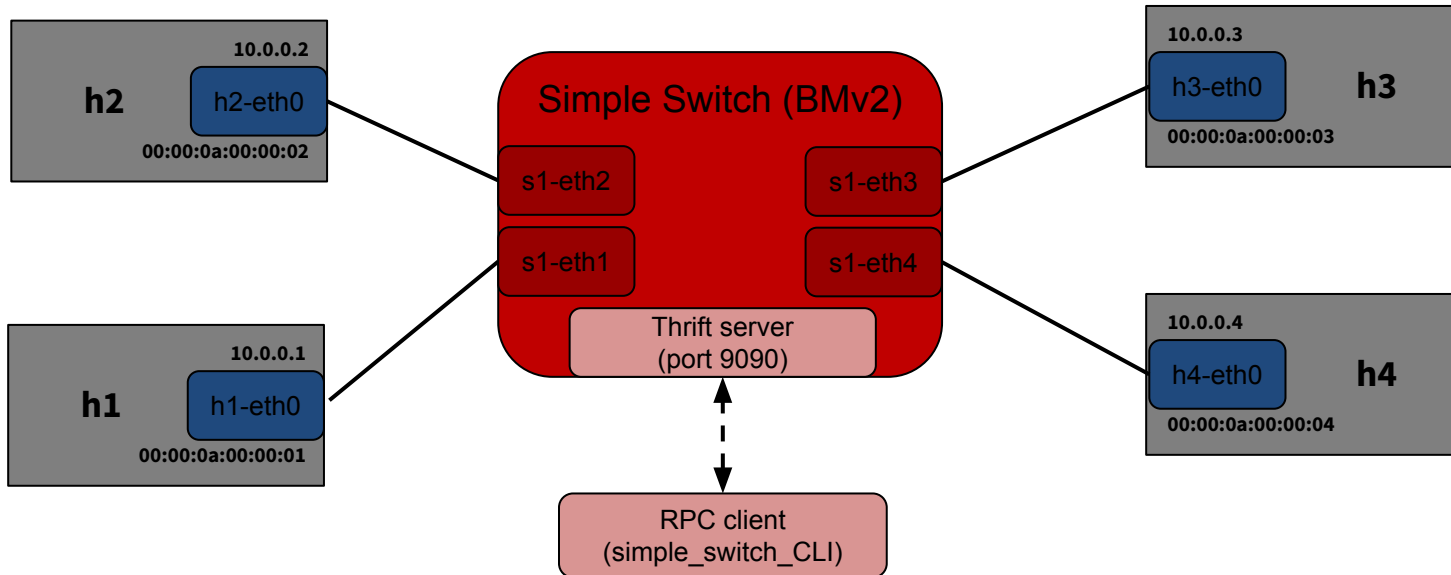
L2 Switching - Solution

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {

    action drop() {
        mark_to_drop(standard_metadata);
    }

    action eth_forward(bit<9> out_port) {
        standard_metadata.egress_spec = out_port;
    }
```

```
    table out_iface {
        key = {
            hdr.ethernet.dstAddr: exact;
        }
        actions = {
            drop;
            eth_forward;
        }
        size = 8;
        default_action = drop;
    }

    apply {
        out_iface.apply();
    }
}
```

- To populate the table...

```
simple_switch_CLI --thrift-port 9090
> table_add out_iface eth_forward 00:00:0a:00:00:01 => 1
> ...
> table_add out_iface eth_forward 00:00:0a:00:00:04 => 4
```

L2 Switching with Access Control List (ACL)

- Switching table provided through the control plane
- Table populated through the simple_switch_CLI (gRPC client)
- Allow forwarding only from mac address included in ACL
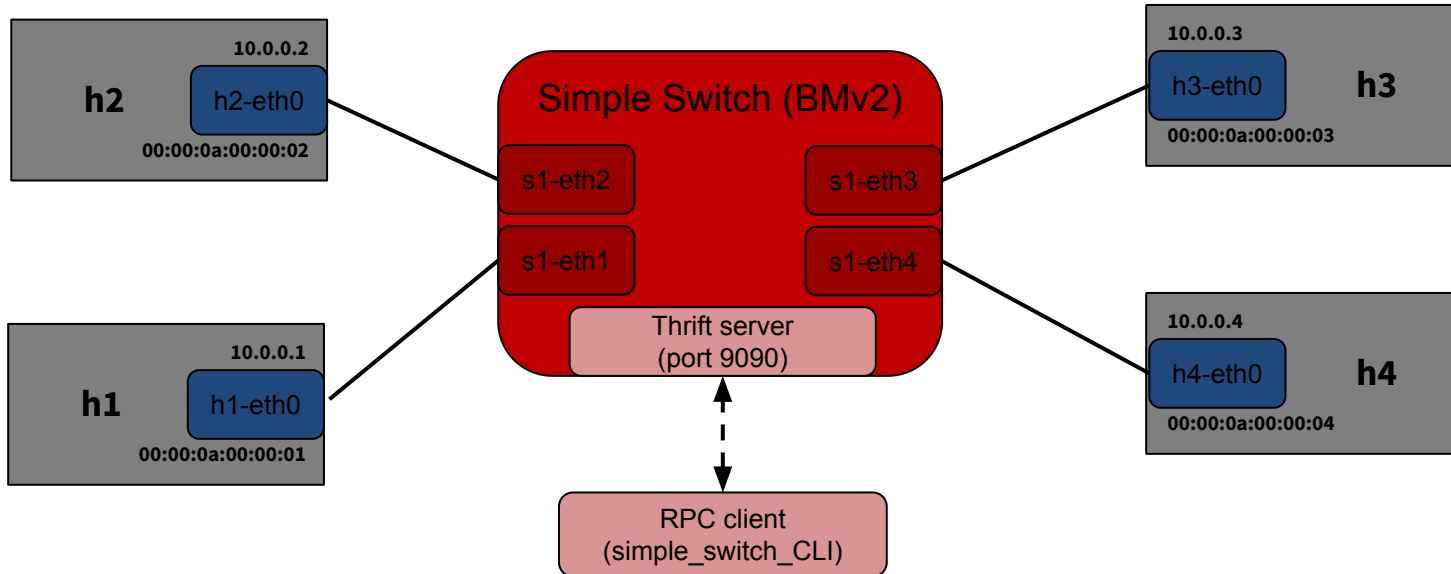
L2 Switching with ACL - Solution

```
control MyEgress(inout headers hdr,
                 inout metadata meta,
                 inout standard_metadata_t standard_metadata) {

    action block() {
        mark_to_drop(standard_metadata);
    }
    table acl {
    key = {
            hdr.ethernet.srcAddr: exact;
    }
    actions = {
                block;
                NoAction;
    }
    default_action = block();
    size = 8;
    }
    apply {
        acl.apply();
    }
}
```

Add to the previous example this egress control

● To populate the table...
```
simple_switch_CLI --thrift-port 9090
> table_add acl NoAction 00:00:0a:00:00:02 =>
> table_add acl NoAction 00:00:0a:00:00:03 =>
```

L3 Forwarding

- Forwarding table provided via the control plane
- Table populated through the simple_switch_CLI (gRPC client)
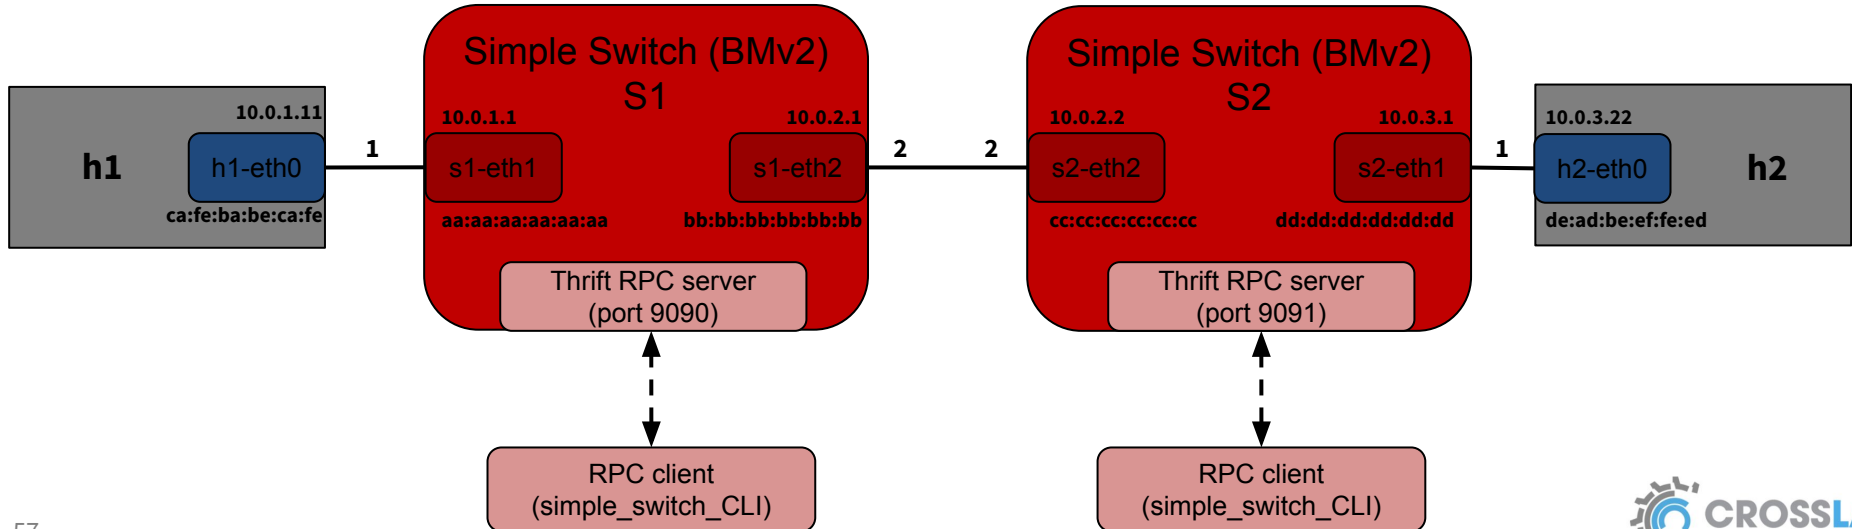
L3-forwarding - Headers

```
const bit<16> TYPE_IPV4 = 0x800;

typedef bit<9>  egressSpec_t;
typedef bit<48> macAddr_t;
typedef bit<32> ip4Addr_t;

header ethernet_t {
    macAddr_t dstAddr;
    macAddr_t srcAddr;
    bit<16>   etherType;
}

header ipv4_t {
    bit<4>     version;
    bit<4>     ihl;
    bit<8>     diffserv;
    bit<16>    totalLen;
    bit<16>    identification;
    bit<3>     flags;
    bit<13>    fragOffset;
    bit<8>     ttl;
    bit<8>     protocol;
    bit<16>    hdrChecksum;
    ip4Addr_t srcAddr;
    ip4Addr_t dstAddr;
}
```

```
struct metadata {
    /* empty */
}

struct headers {
    ethernet_t    ethernet;
    ipv4_t        ipv4;
}
```

L3-forwarding - Parser

```
parser MyParser(packet_in packet,
                out headers hdr,
                inout metadata meta,
                inout standard_metadata_t standard_metadata) {

    state start {

        packet.extract(hdr.ethernet);
        transition select(hdr.ethernet.etherType){
            TYPE_IPV4: ipv4;  //TYPE_PIV4 = 0x0800
            default: accept;
        }
    }

    state ipv4 {
        packet.extract(hdr.ipv4);
        transition accept;
    }

}
```

L3-forwarding - Control flow

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {

    action drop() {
        mark_to_drop(standard_metadata);
    }

    action ipv4_forward(macAddr_t d_mac, egressSpec_t port){

        hdr.ethernet.dstAddr = d_mac;

        standard_metadata.egress_spec = port;

        hdr.ipv4.ttl = hdr.ipv4.ttl - 1;
    }
```

```
table ipv4_lpm {
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    actions = {
        ipv4_forward;
        drop;
        NoAction;
    }
    size = 1024;
    default_action = NoAction;
}


apply {
    if ( hdr.ipv4.isValid() ) {
        ipv4_lpm.apply();
    }
}
}
```

L3-forwarding - Egress processing

```
control MyEgress(inout headers hdr,
                 inout metadata meta,
                 inout standard_metadata_t standard_metadata) {

    action drop() {
        mark_to_drop(standard_metadata);
    }

    action set_smac(macAddr_t mac) {
        hdr.ethernet.srcAddr = mac;
    }

    table s_mac {
        key = {
            standard_metadata.egress_port: exact;
        }
        actions = {
            set_smac;
            drop;
            NoAction;
        }
        size = 16;
        default_action = NoAction;
    }
```

```
    apply {
        s_mac.apply();
    }
}
```

L3 forwarding - Checksum and deparser

```
control MyComputeChecksum(inout headers hdr, inout metadata meta)
{
    apply {
        update_checksum(
        hdr.ipv4.isValid(),
         { hdr.ipv4.version,
          hdr.ipv4.ihl,
            hdr.ipv4.diffserv,
            hdr.ipv4.totalLen,
            hdr.ipv4.identification,
            hdr.ipv4.flags,
            hdr.ipv4.fragOffset,
            hdr.ipv4.ttl,
            hdr.ipv4.protocol,
            hdr.ipv4.srcAddr,
            hdr.ipv4.dstAddr },
        hdr.ipv4.hdrChecksum,
        HashAlgorithm.csum16);
    }
}
```

```
control MyDeparser(packet_out packet, in headers hdr) {
    apply {

        packet.emit(hdr.ethernet);
        packet.emit(hdr.ipv4);

    }
}
```

```
//switch architecture
V1Switch(
MyParser(),
MyVerifyChecksum(),
MyIngress(),
MyEgress(),
MyComputeChecksum(),
MyDeparser()
) main;
```

Innovation for industry 4.0

L3 Forwarding: populating the tables

```
> simple_switch_CLI --thrift-port 9090 < s1-commands.txt

> simple_switch_CLI --thrift-port 9091 < s2-commands.txt
```

```
# s1-commands.txt

table_clear ipv4_lpm
table_add ipv4_lpm ipv4_forward 10.0.1.11/32 => ca:fe:ba:be:ca:fe 1
table_add ipv4_lpm ipv4_forward 10.0.3.0/24  => cc:cc:cc:cc:cc:cc 2

table_clear s_mac
table_add s_mac set_smac 1 => aa:aa:aa:aa:aa:aa
table_add s_mac set_smac 2 => bb:bb:bb:bb:bb:bb
```

```
# s2-commands.txt

table_clear ipv4_lpm
table_add ipv4_lpm ipv4_forward 10.0.1.0/24  => bb:bb:bb:bb:bb:bb 2
table_add ipv4_lpm ipv4_forward 10.0.3.22/32 => de:ad:be:ef:fe:ed 1

table_clear s_mac
table_add s_mac set_smac 1 => dd:dd:dd:dd:dd:dd
table_add s_mac set_smac 2 => cc:cc:cc:cc:cc:cc
```

IP traffic probe and filter

- The switch normally forward frames among h1, h2 and h3
- in addition, it "mirrors" (clone) every IP packet to the probe
  - Traffic to the probe is further filtered by *protocol* (in the IP header)
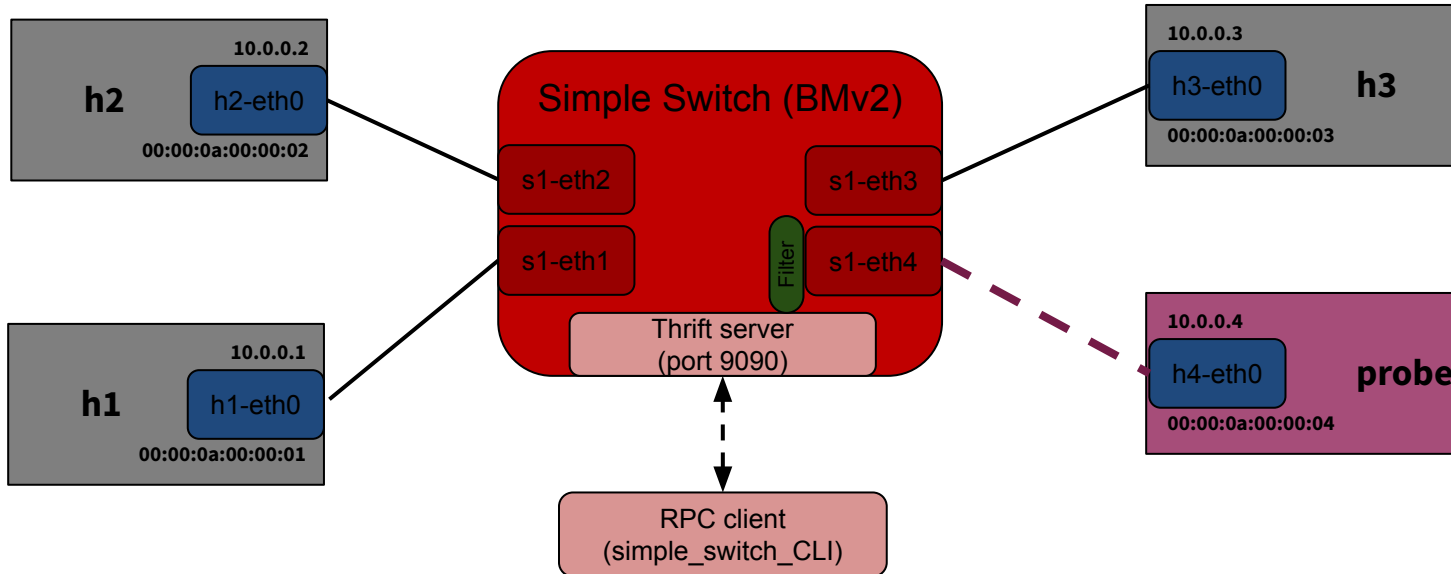
IP probe - Headers

```
const bit<16> TYPE_IPV4 = 0x800;

typedef bit<48> macAddr_t;
typedef bit<32> ip4Addr_t;

header ethernet_t {
    macAddr_t dstAddr;
    macAddr_t srcAddr;
    bit<16>   etherType;
}

header ipv4_t {
    bit<4>    version;
    bit<4>    ihl;
    bit<8>    diffserv;
    bit<16>   totalLen;
    bit<16>   identification;
    bit<3>    flags;
    bit<13>   fragOffset;
    bit<8>    ttl;
    bit<8>    protocol;
    bit<16>   hdrChecksum;
    ip4Addr_t srcAddr;
    ip4Addr_t dstAddr;
}
```

```
struct metadata {
    /* empty */
}

struct headers {
    ethernet_t    ethernet;
    ipv4_t        ipv4;
}
```

IP probe - Parser

```
parser MyParser(packet_in packet,
                out headers hdr,
                inout metadata meta,
                inout standard_metadata_t standard_metadata) {

    state start {

        packet.extract(hdr.ethernet);
        transition select(hdr.ethernet.etherType){
            TYPE_IPV4: ipv4;  //TYPE_PIV4 = 0x0800
            default: accept;
        }
    }

    state ipv4 {
        packet.extract(hdr.ipv4);
        transition accept;
    }

}
```

IP probe - Control flow

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {
    action drop() {
        mark_to_drop(standard_metadata);
    }
    action eth_forward(bit<9> out_port) {
        standard_metadata.egress_spec = out_port;
    }

    table out_iface {
        key = {
            hdr.ethernet.dstAddr: exact;
        }
        actions = {
            drop;
            eth_forward;
        }
        size = 8;
        default_action = drop;
    }
```

```
    action pkt_clone(bit<32> mirror_id) {
        clone(CloneType.I2E, mirror_id);
    }

    table monitor {
        key = {
            hdr.ipv4.protocol: exact;
        }
        actions = {
            NoAction;
            pkt_clone;
        }
        size = 8;
        default_action = NoAction;
    }


    apply {
        if ( hdr.ipv4.isValid() ) {
            monitor.apply();
        }

        out_iface.apply();
    }
}
```

IP probe: populating the tables

```
> simple_switch_CLI --thrift-port 9090 < s1-commands.txt
```

```
# s1-commands.txt

table_add out_iface eth_forward 00:00:0a:00:00:01 => 1
table_add out_iface eth_forward 00:00:0a:00:00:02 => 2
table_add out_iface eth_forward 00:00:0a:00:00:03 => 3

mirroring_add 100 4
table_add monitor pkt_clone 1 => 100
table_add monitor pkt_clone 6 => 100
```

Part II

Giuseppe Lettieri

End host data plane programmability